

Introduction to R

Table of contents

Welcome	4
Source code	4
Authors and copyright	5
1 Starting out in R	6
1.1 Variables	8
1.2 Saving code in an R script	9
1.3 Vectors	9
1.4 Types of vector	10
1.5 Indexing vectors	11
Challenge: indexing	12
1.6 Sequences	12
1.7 Functions	13
Challenge: using functions	15
2 Data frames	16
2.1 Setting up	16
2.2 Loading data	17
2.3 Exploring	19
2.4 Indexing data frames	20
2.5 Columns are vectors	23
Quiz: reading R code	24
2.6 Logical indexing	24
Challenge: logical indexing	26
2.6.1 A <code>dplyr</code> shorthand	27
2.7 Factors	27
2.8 Readability vs tidyness	30
2.9 Sorting	31
2.10 Joining data frames	32
Quiz	32
2.11 Further reading	33
3 Plotting with ggplot2	34
3.1 Elements of a ggplot	34
Challenge: make a ggplot	36

3.2	Further geoms	36
3.3	Highlighting subsets	39
	Discussion: look at the plot	40
3.4	Fine-tuning a plot	40
	Challenge: refine your ggplot	42
3.5	Faceting	42
	Challenge: facet your ggplot	43
3.6	Saving ggplots	43
4	Summarizing data	45
4.1	Summary functions	45
4.2	Missing values	46
4.3	Grouped summaries	47
	Challenge: summarizing	48
4.4	t-test	50
5	Thinking in R	52
5.1	Lists	53
5.2	Other types not covered here	54
5.3	Programming	54
6	Next steps	56
6.1	Deepen your understanding	56
6.2	Expand your vocabulary	56
6.3	Find packages for specific data types	56
6.4	Some pointers on models and statistics	57

Welcome

This is workshop material for the “Introduction to R” workshop given by the Monash Genomics and Bioinformatics Platform¹. In the past it has been presented as part of the Data Fluency program at Monash. Our teaching style is based on the style of The Carpentries².

- Slideshow³
- PDF version for printing⁴
- ZIP of data files used in this workshop⁵

During the workshop we will be using Posit Cloud to use R over the web:

- Posit Cloud⁶

You can also install R on your own computer. There are two things to download and install:

- Download R⁷
- Download RStudio⁸

R is the language itself. RStudio provides a convenient environment in which to use R, either on your local computer or on a server.

Source code

This book was created in R using quarto⁹!

- GitHub page¹⁰

¹<https://www.monash.edu/researchinfrastructure/mgbp>

²<https://carpentries.org/>

³[slides.html](#)

⁴[r-intro-2.pdf](#)

⁵[r-intro.zip](#)

⁶<https://posit.cloud/>

⁷<https://cran.rstudio.com/>

⁸<https://www.rstudio.com/products/rstudio/download/>

⁹<https://quarto.org/>

¹⁰<https://github.com/MonashDataFluency/r-intro-2>


Authors and copyright

This course is developed for the Monash Genomics and Bioinformatics Platform by Paul Harrison.



This work is licensed under a CC BY-4: Creative Commons Attribution 4.0 International License¹¹. The attribution is “Monash Genomics and Bioinformatics Platform” if copying or modifying these notes.

Data files are derived from Gapminder, which has a CC BY-4 license. The attribution is “Free data from www.gapminder.org”. The data is given here in a form designed to teach various points about the R language. Refer to the Gapminder site¹² for the original form of the data if using it for other uses.



MONASH GENOMICS AND BIOINFORMATICS PLATFORM



¹¹<http://creativecommons.org/licenses/by/4.0/>

¹²<https://www.gapminder.org>

1 Starting out in R

R is both a programming language and an interactive environment for data exploration and statistics. Today we will be concentrating on R as an *interactive environment*.

Working with R is primarily text-based. The basic mode of use for R is that the user types in a command in the R language and presses enter, and then R computes and displays the result.

We will be working in RStudio¹. The easiest way to get started is to go to Posit Cloud² and create a new project. Monash staff and students can log in using their Monash Google account.

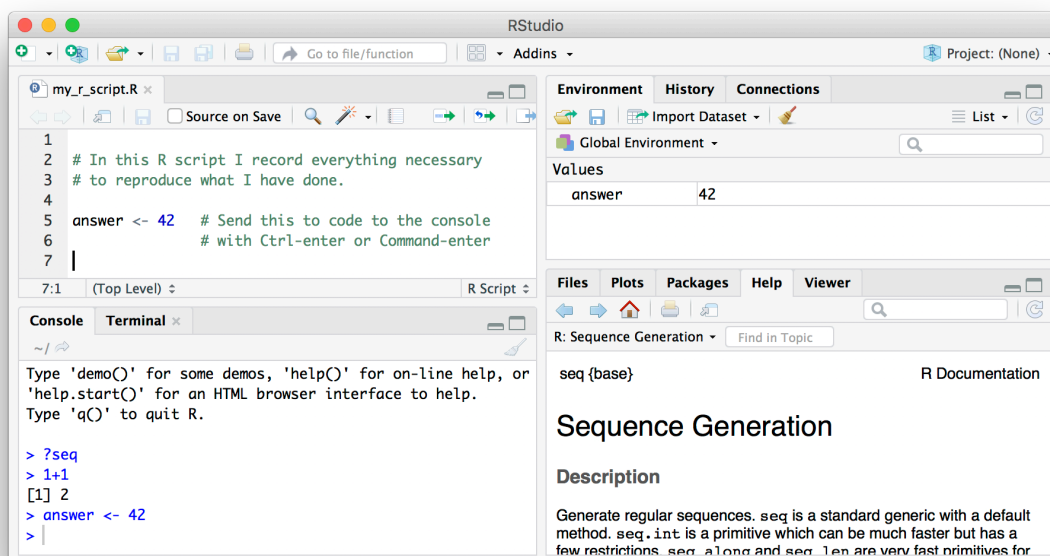
If you are using RStudio directly on your own laptop, we also recommend creating a project for this workshop. After opening the desktop version of RStudio, go to the “File” menu and select “New Project...”. This will keep any files you create in their own folder.

The main way of working with R is the *console*, where you enter commands and view results. RStudio surrounds this with various conveniences. In addition to the console panel, RStudio provides panels containing:

- A text editor, where R commands can be recorded for future reference.
- A history of commands that have been typed on the console.
- An “environment” pane with a list of *variables*, which contain values that R has been told to save from previous commands.
- A file manager.
- Help on the functions available in R.
- A panel to show plots.

¹<https://www.rstudio.com/products/rstudio/download/>

²<https://rstudio.cloud/>



Open RStudio, click on the “Console” pane, type `1+1` and press enter. R displays the result of the calculation. In this document, we will show such an interaction with R as below.

```
1+1
```

```
[1] 2
```

`+` is called an operator. R has the operators you would expect for mathematics: `+` `-` `*` `/` `^`. It also has further operators that do other things.

`*` has higher precedence than `+`. We can use brackets if necessary `()`. Try `1+2*3` and `(1+2)*3`.

Spaces can be used to make code easier to read.

We can compare with `==` `<` `>` `<=` `>=` `!=`. This produces a *logical* value, `TRUE` or `FALSE`. Note the double equals, `==`, for equality comparison.

```
2 * 2 == 4
```

```
[1] TRUE
```

There are also character strings such as `"string"`. A character string must be surrounded by either single or double quotes.

1.1 Variables

A variable is a name for a value. We can create a new variable by assigning a value to it using `<-`.

```
width <- 5
```

RStudio helpfully shows us the variable in the “Environment” pane. We can also print it by typing the name of the variable and hitting enter. In general, R will print to the console any object returned by a function or operation *unless* we assign it to a variable.

```
width
```

```
[1] 5
```

Examples of valid variables names: `hello`, `subject_id`, `subject.ID`, `x42`. Spaces aren’t ok *inside* variable names. Dots (.) are ok in R, unlike in many other languages. Numbers are ok, except as the first character. Punctuation is not allowed, with two exceptions: `_` and `..`

We can do arithmetic with the variable:

```
# Area of a square  
width * width
```

```
[1] 25
```

and even save the result in another variable:

```
# Save area in "area" variable  
area <- width * width
```

We can also change a variable’s value by assigning it a new value:

```
width <- 10  
width
```

```
[1] 10
```



```
area
```

```
[1] 25
```

Notice that the value of `area` we calculated earlier hasn't been updated. Assigning a new value to one variable does not change the values of other variables. This is different to a spreadsheet, but usual for programming languages.

1.2 Saving code in an R script

Once we've created a few variables, it becomes important to record how they were calculated so we can reproduce them later.

The usual workflow is to save your code in an R script (".R file"). Go to "File/New File/R Script" to create a new R script. Code in your R script can be sent to the console by selecting it or placing the cursor on the correct line, and then pressing **Control-Enter** (**Command-Enter** on a Mac).

Tip

Add comments to code, using lines starting with the `#` character. This makes it easier for others to follow what the code is doing (and also for us the next time we come back to it).

1.3 Vectors

A *vector* of numbers is a collection of numbers. "Vector" means different things in different fields (mathematics, geometry, biology), but in R it is a fancy name for a collection of numbers. We call the individual numbers *elements* of the vector.

We can make vectors with `c()`, for example `c(1,2,3)`. `c` means "combine". R is obsessed with vectors, in R even single numbers are vectors of length one. Many things that can be done with a single number can also be done with a vector. For example arithmetic can be done on vectors as it can be on single numbers.

```
myvec <- c(10,20,30,40,50)
myvec
```

```
[1] 10 20 30 40 50
```

```
myvec + 1
```

```
[1] 11 21 31 41 51
```

```
myvec + myvec
```

```
[1] 20 40 60 80 100
```

```
length(myvec)
```

```
[1] 5
```

```
c(60, myvec)
```

```
[1] 60 10 20 30 40 50
```

```
c(myvec, myvec)
```

```
[1] 10 20 30 40 50 10 20 30 40 50
```

When we talk about the length of a vector, we are talking about the number of numbers in the vector.

1.4 Types of vector

We will also encounter vectors of character strings, for example "hello" or `c("hello", "world")`. Also we will encounter “logical” vectors, which contain `TRUE` and `FALSE` values. R also has “factors”, which are categorical vectors, and behave much like character vectors (think the factors in an experiment).

Because vectors can only contain one type of thing, when you mix different types R will choose a lowest common denominator type of vector, a type that can contain everything we are trying to put in it. A different language might stop with an error, but R tries to soldier on as best it can. A number can be represented as a character string, but a character string can not be represented as a number, so when we try to put both in the same vector R converts everything to a character string.

```
c("hello", 1, TRUE)
```

```
[1] "hello" "1"      "TRUE"
```

1.5 Indexing vectors

Access elements of a vector with `[]`, for example `myvec[1]` to get the first element. You can also assign to a specific element of a vector.

```
myvec[1]
```

```
[1] 10
```

```
myvec[2]
```

```
[1] 20
```

```
myvec[2] <- 5  
myvec
```

```
[1] 10  5 30 40 50
```

Can we use a vector to index another vector? Yes!

```
myind <- c(4,3,2)  
myvec[myind]
```

```
[1] 40 30  5
```

We could equivalently have written:

```
myvec[c(4,3,2)]
```

```
[1] 40 30  5
```

Challenge: indexing

We can create and index character vectors as well. A cafe is using R to create their menu.

```
menu <- c("spam", "eggs", "beans", "bacon", "sausage")
```

1. What does `menu[-3]` produce? Based on what you find, use indexing to create a version of `menu` without "spam".
2. Use indexing to create a vector containing spam, eggs, sausage, spam, and spam.
3. Add a new item, "lobster", to `menu`, and store the result in a variable called `new_menu`.

1.6 Sequences

Another way to create a vector is with `::`

```
1:10
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

This can be useful when combined with indexing:

```
menu[1:4]
```


```
[1] "spam" "eggs" "beans" "bacon"
```

Sequences are useful for other things, such as a starting point for calculations:

```
x <- 1:10  
x*x
```

```
[1] 1 4 9 16 25 36 49 64 81 100
```

```
plot(x, x*x)
```



1.7 Functions

Functions are the things that do all the work for us in R: calculate, manipulate data, read and write to files, produce plots. R has many built in functions and we will also be loading more specialized functions from “packages”.

We’ve already seen several functions: `c()`, `length()`, and `plot()`. Let’s now have a look at `sum()`.

```
sum(myvec)
```

```
[1] 135
```

We *called* the function `sum` with the *argument* `myvec`, and it *returned* the value 135. We can get help on how to use `sum` with:

```
?sum
```

Some functions take more than one argument. Let’s look at the function `rep`, which means “repeat”, and which can take a variety of different arguments. In the simplest case, it takes a value and the number of times to repeat that value.

```
rep(42, 10)
```

```
[1] 42 42 42 42 42 42 42 42 42 42
```

As with many functions in R—which is obsessed with vectors—the thing to be repeated can be a vector with multiple elements.

```
rep(c(1,2,3), 10)
```

```
[1] 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3
```

So far we have used *positional* arguments, where R determines which argument is which by the order in which they are given. We can also give arguments by *name*. For example, the above is equivalent to

```
rep(c(1,2,3), times=10)
```

```
[1] 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3
```

```
rep(x=c(1,2,3), 10)
```

```
[1] 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3
```

```
rep(times=10, x=c(1,2,3))
```

```
[1] 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3
```

Arguments can have default values, and a function may have many different possible arguments that make it do obscure things. For example, `rep` can also take an argument `each=`. It's typical for a function to be invoked with some number of positional arguments, which are always given, plus some less commonly used arguments, typically given by name.

```
rep(c(1,2,3), each=3)
```

```
[1] 1 1 1 2 2 2 3 3 3
```

```
rep(c(1,2,3), each=3, times=5)
```

```
[1] 1 1 1 2 2 2 3 3 3 1 1 1 2 2 2 3 3 3 1 1 1 2 2 2 3 3 3 1 1 1 2 2 2 3 3 3 1 1 1  
[39] 1 2 2 2 3 3 3
```

Challenge: using functions

1. Use `sum` to sum from 1 to 10,000.
2. Look at the documentation for the `seq` function. What does `seq` do? Give an example of using `seq` with either the `by` or `length.out` argument.

2 Data frames

Data frame is R's name for tabular data. We generally want each row in a data frame to represent a unit of observation, and each column to contain a different type of information about the units of observation. Tabular data in this form is called “tidy data”¹.

Today we will be using a collection of modern packages collectively known as the Tidyverse². R and its predecessor S have a history dating back to 1976. The Tidyverse fixes some dubious design decisions baked into “base R”, including having its own slightly improved form of data frame, which is called a *tibble*. Sticking to the Tidyverse where possible is generally safer, Tidyverse packages are more willing to generate errors rather than ignore problems.

2.1 Setting up

Our first step is to download the files we need and to install the Tidyverse. This is the one step where we ask you to copy and paste some code:

```
# Download files for this workshop
download.file(
  "https://monashdatafluency.github.io/r-intro-2/r-intro.zip",
  destfile="r-intro.zip")
unzip("r-intro.zip")

# Install Tidyverse
install.packages("tidyverse")
```

If you run into problems installing all of the `tidyverse` you may have more success installing individual packages:

```
install.packages(c("dplyr", "readr", "tidyr", "ggplot2"))
```

We need to load the `tidyverse` package in order to use it.

¹<http://vita.had.co.nz/papers/tidy-data.html>

²<https://www.tidyverse.org/>


```
library(tidyverse)

# OR
library(dplyr)
library(readr)
library(tidyr)
library(ggplot2)
```

The `tidyverse` package loads various other packages, setting up a modern R environment. In this section we will be using functions from the `dplyr`, `readr` and `tidyr` packages.

R is a language with mini-languages within it that solve specific problem domains. `dplyr` is such a mini-language, a set of “verbs” (functions) that work well together. `dplyr`, with the help of `tidyr` for some more complex operations, provides a way to perform most manipulations on a data frame that you might need.

2.2 Loading data

We will use the `read_csv` function from `readr` to load a data set. (See also `read.csv` in base R.) CSV stands for Comma Separated Values, and is a text format used to store tabular data. The first few lines of the file we are loading are shown below. Conventionally the first line contains column headings.

```
name,region,oezd,g77,lat,long,income2017
Afghanistan,asia,FALSE,TRUE,33,66,low
Albania,europe,FALSE,FALSE,41,20,upper_mid
Algeria,africa,FALSE,TRUE,28,3,upper_mid
Andorra,europe,FALSE,FALSE,42.50779,1.52109,high
Angola,africa,FALSE,TRUE,-12.5,18.5,lower_mid
```

```
geo <- read_csv("r-intro/geo.csv")
```

```
Rows: 196 Columns: 7
```

```
-- Column specification -----
Delimiter: ","
chr (3): name, region, income2017
dbl (2): lat, long
lgl (2): oezd, g77
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
# A tibble: 196 x 7
  name          region oecd  g77    lat  long income2017
  <chr>         <chr> <lgl> <lgl> <dbl> <dbl> <chr>
1 Australia    asia   TRUE  FALSE -25    135  high
2 Brunei       asia   FALSE TRUE   4.5   115. high
3 Cambodia     asia   FALSE TRUE   13    105  lower_mid
4 China        asia   FALSE TRUE   35    105  upper_mid
5 Fiji         asia   FALSE TRUE  -18    178  upper_mid
6 Hong Kong, China asia   FALSE FALSE  22.3  114. high
7 Indonesia    asia   FALSE TRUE   -5    120  lower_mid
8 Japan        asia   TRUE  FALSE  35.7  140. high
9 Kiribati     asia   FALSE FALSE   1.42  173. lower_mid
10 North Korea asia   FALSE TRUE   40    127  low
# i 186 more rows
```

`read_csv` has guessed the type of data each column holds:

- `<chr>` - character strings
- `<dbl>` - numerical values. Technically these are “doubles”, which is a way of storing numbers with 15 digits precision.
- `<lgl>` - logical values, TRUE or FALSE.

We will also encounter:

- `<int>` - integers, a fancy name for whole numbers.
- `<fct>` - factors, categorical data. We will get to this shortly.

You can also see this data frame referring to itself as “a tibble”. This is the Tidyverse’s improved form of data frame. Tibbles present themselves more conveniently than base R data frames. Base R data frames don’t show the type of each column, and output every row when you try to view them.

Tip

A data frame can also be created from vectors, with the `tibble` function. (See also `data.frame` in base R.) For example:

```
tibble(foo=c(10,20,30), bar=c("a","b","c"))
```

```
# A tibble: 3 x 2
  foo bar
```

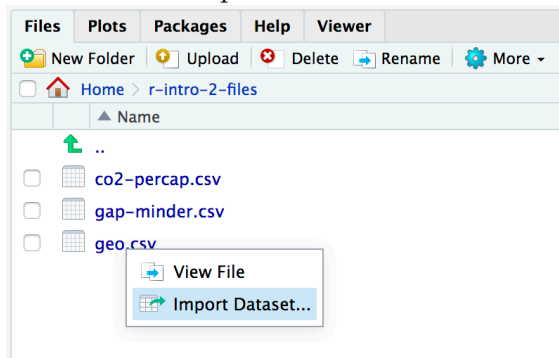
```
<dbl> <chr>
1      10 a
2      20 b
3      30 c
```

The argument names become column names in the data frame.

💡 Tip

The *path* to the file on our server is "`r-intro/geo.csv`". This says, starting from your working directory, look in the directory `r-intro` for the file `geo.csv`. The steps in the path are separated by `/`. Your working directory is shown at the top of the console pane. The path needed might be different on your own computer, depending where you downloaded the file.

One way to work out the correct path is to find the file in the file browser pane, click on it and select "Import Dataset...".



2.3 Exploring

The `View` function gives us a spreadsheet-like view of the data frame.

```
View(geo)
```

`print` with the `n` argument can be used to show more than the first 10 rows on the console.

```
print(geo, n=200)
```

We can extract details of the data frame with further functions:

```
nrow(geo)
```

```
[1] 196
```

```
ncol(geo)
```

```
[1] 7
```

```
colnames(geo)
```

```
[1] "name"      "region"    "oecd"      "g77"       "lat"
[6] "long"      "income2017"
```

```
summary(geo)
```

name	region	oecd	g77
Length:196	Length:196	Mode :logical	Mode :logical
Class :character	Class :character	FALSE:165	FALSE:65
Mode :character	Mode :character	TRUE :31	TRUE :131

lat	long	income2017
Min. : -42.00	Min. : -175.000	Length:196
1st Qu.: 4.00	1st Qu.: -5.625	Class :character
Median : 17.42	Median : 21.875	Mode :character
Mean : 19.03	Mean : 23.004	
3rd Qu.: 39.82	3rd Qu.: 51.892	
Max. : 65.00	Max. : 179.145	

2.4 Indexing data frames

Data frames can be subset using `[row,column]` syntax.

```
geo[4,2]
```

```
# A tibble: 1 x 1
  region
  <chr>
1 asia
```

Note that while this is a single value, it is still wrapped in a data frame. (This is a behaviour specific to Tidyverse data frames.) More on this in a moment.

Columns can be given by name.

```
geo[4,"region"]
```

```
# A tibble: 1 x 1
  region
  <chr>
1 asia
```

The column or row may be omitted, thereby retrieving the entire row or column.

```
geo[4,]
```

```
# A tibble: 1 x 7
  name region oecd g77 lat long income2017
  <chr> <chr> <lgl> <lgl> <dbl> <dbl> <chr>
1 China asia FALSE TRUE 35 105 upper_mid
```

```
geo[, "region"]
```

```
# A tibble: 196 x 1
  region
  <chr>
1 asia
2 asia
3 asia
4 asia
5 asia
6 asia
7 asia
8 asia
9 asia
10 asia
# i 186 more rows
```

Multiple rows or columns may be retrieved using a vector.

```
rows_wanted <- c(1,3,5)
geo[rows_wanted,]
```

```
# A tibble: 3 x 7
  name      region oecd g77    lat  long income2017
  <chr>    <chr> <lgl> <lgl> <dbl> <dbl> <chr>
1 Australia asia  TRUE FALSE  -25  135 high
2 Cambodia asia  FALSE TRUE   13  105 lower_mid
3 Fiji     asia  FALSE TRUE  -18  178 upper_mid
```

Vector indexing can also be written on a single line.

```
geo[c(1,3,5),]
```

```
# A tibble: 3 x 7
  name      region oecd g77    lat  long income2017
  <chr>    <chr> <lgl> <lgl> <dbl> <dbl> <chr>
1 Australia asia  TRUE FALSE  -25  135 high
2 Cambodia asia  FALSE TRUE   13  105 lower_mid
3 Fiji     asia  FALSE TRUE  -18  178 upper_mid
```

```
geo[1:7,]
```

```
# A tibble: 7 x 7
  name      region oecd g77    lat  long income2017
  <chr>    <chr> <lgl> <lgl> <dbl> <dbl> <chr>
1 Australia asia  TRUE FALSE  -25  135 high
2 Brunei   asia  FALSE TRUE   4.5 115. high
3 Cambodia asia  FALSE TRUE   13  105 lower_mid
4 China    asia  FALSE TRUE   35  105 upper_mid
5 Fiji     asia  FALSE TRUE  -18  178 upper_mid
6 Hong Kong, China asia FALSE FALSE 22.3 114. high
7 Indonesia asia  FALSE TRUE  -5   120 lower_mid
```

2.5 Columns are vectors

Ok, so how do we actually get data out of a data frame?

Under the hood, a data frame is a list of column vectors. We can use `$` to retrieve columns. Occasionally it is also useful to use `[[]]` to retrieve columns, for example if the column name we want is stored in a variable.

```
head( geo$region )
```

```
[1] "asia" "asia" "asia" "asia" "asia" "asia"
```

```
head( geo[["region"]] )
```

```
[1] "asia" "asia" "asia" "asia" "asia" "asia"
```


To get the “region” value of the 4th row as above, but unwrapped, we can use:

```
geo$region[4]
```

```
[1] "asia"
```

For example, to plot the longitudes and latitudes we could use:

```
plot(geo$long, geo$lat)
```



Quiz: reading R code

You encounter some wild R code, and aren't sure what it does. Based on R syntax you've encountered so far, what roles are the different names in this code playing?

```
highest <- geo$name[ head(order(geo$lat, decreasing=TRUE), n=10) ]
```

Find all examples of:

- A. The name of a variable to store a value to.
- B. The name of a variable to retrieve the value from.
- C. The name of a column to get from a data frame.
- D. The name of a function to call.
- E. The name of an argument to a function call.

See [here³](#) for answers.

2.6 Logical indexing

A method of indexing that we haven't discussed yet is logical indexing. Instead of specifying the row number or numbers that we want, we can give a logical vector which is **TRUE** for the rows we want and **FALSE** otherwise. This can also be used with vectors.

We will first do this in a slightly verbose way in order to understand it, then learn a more concise way to do this using the **dplyr** package.

Southern countries have latitude less than zero.

```
is_southern <- geo$lat < 0
```

```
head(is_southern)
```

```
[1]  TRUE FALSE FALSE FALSE  TRUE FALSE
```

```
sum(is_southern)
```

```
[1] 40
```

³[answers/answers-reading-r-code.html](#)

`sum` treats TRUE as 1 and FALSE as 0, so it tells us the number of TRUE elements in the vector.

We can use this logical vector to get the southern countries from `geo`:

```
geo[is_southern,]
```

```
# A tibble: 40 x 7
  name          region oecd  g77      lat  long income2017
  <chr>         <chr> <lgl> <lgl>   <dbl> <dbl> <chr>
1 Australia    asia   TRUE  FALSE -25     135  high
2 Fiji         asia   FALSE TRUE  -18     178  upper_mid
3 Indonesia    asia   FALSE TRUE   -5     120  lower_mid
4 Nauru         asia   FALSE FALSE -0.517  167.  upper_mid
5 New Zealand  asia   TRUE  FALSE -42     174  high
6 Papua New Guinea asia   FALSE TRUE   -6     147  lower_mid
7 Samoa        asia   FALSE TRUE -13.8   -172.  upper_mid
8 Solomon Islands asia   FALSE TRUE   -8     159  lower_mid
9 Timor-Leste  asia   FALSE TRUE  -8.83   126.  lower_mid
10 Tonga        asia   FALSE TRUE  -20    -175  upper_mid
# i 30 more rows
```

Comparison operators available are:

- `x == y` – “equal to”
- `x != y` – “not equal to”
- `x < y` – “less than”
- `x > y` – “greater than”
- `x <= y` – “less than or equal to”
- `x >= y` – “greater than or equal to”

More complicated conditions can be constructed using logical operators:

- `a & b` – “and”, TRUE only if both `a` and `b` are TRUE.
- `a | b` – “or”, TRUE if either `a` or `b` or both are TRUE.
- `! a` – “not”, TRUE if `a` is FALSE, and FALSE if `a` is TRUE.

The `oecd` column of `geo` tells which countries are in the Organisation for Economic Co-operation and Development, and the `g77` column tells which countries are in the Group of 77 (an alliance of developing nations). We could see which OECD countries are in the southern hemisphere with:

```
southern_oecd <- is_southern & geo$oecd
geo[southern_oecd,]
```

```
# A tibble: 3 x 7
  name      region  oecd  g77    lat  long income2017
  <chr>    <chr>   <lgl> <lgl> <dbl> <dbl> <chr>
1 Australia asia    TRUE  FALSE -25   135   high
2 New Zealand asia    TRUE  FALSE -42   174   high
3 Chile    americas TRUE   TRUE  -33.5 -70.6 high
```

`is_southern` seems like it should be kept within our `geo` data frame for future use. We can add it as a new column of the data frame with:

```
geo$southern <- is_southern
geo
```

```
# A tibble: 196 x 8
  name      region  oecd  g77    lat  long income2017 southern
  <chr>    <chr>   <lgl> <lgl> <dbl> <dbl> <chr>    <lgl>
1 Australia asia    TRUE  FALSE -25   135   high     TRUE
2 Brunei    asia    FALSE TRUE   4.5  115.   high     FALSE
3 Cambodia asia    FALSE TRUE   13   105   lower_mid FALSE
4 China     asia    FALSE TRUE   35   105   upper_mid FALSE
5 Fiji      asia    FALSE TRUE  -18   178   upper_mid TRUE
6 Hong Kong, China asia    FALSE FALSE 22.3  114.   high     FALSE
7 Indonesia asia    FALSE TRUE   -5   120   lower_mid TRUE
8 Japan     asia    TRUE  FALSE 35.7  140.   high     FALSE
9 Kiribati  asia    FALSE FALSE  1.42 173.   lower_mid FALSE
10 North Korea asia    FALSE TRUE   40   127   low      FALSE
# i 186 more rows
```

Challenge: logical indexing

1. Which country is in both the OECD and the G77?
2. Which countries are in neither the OECD nor the G77?
3. Which countries are in the Americas? These have longitudes between -150 and -40.

2.6.1 A dplyr shorthand

The above method is a little laborious. We have to keep mentioning the name of the data frame, and there is a lot of punctuation to keep track of. `dplyr` provides a slightly magical function called `filter` which lets us write more concisely. For example:

```
filter(gео, lat < 0 & оecd)
```

```
# A tibble: 3 x 8
  name      region оecd g77    lat  long income2017 southern
  <chr>    <chr>   <lgl> <lgl> <dbl> <dbl> <chr>         <lgl>
1 Australia asia    TRUE  FALSE -25   135   high         TRUE
2 New Zealand asia    TRUE  FALSE -42   174   high         TRUE
3 Chile     americas TRUE   TRUE  -33.5 -70.6 high         TRUE
```

In the second argument, we are able to refer to columns of the data frame as though they were variables. The code is beautiful, but also opaque. It's important to understand that under the hood we are creating and combining logical vectors.

2.7 Factors

The `count` function from `dplyr` can help us understand the contents of some of the columns in `geo`. `count` is also *magical*, we can refer to columns of the data frame directly in the arguments to `count`.

```
count(gео, region)
```

```
# A tibble: 4 x 2
  region      n
  <chr>   <int>
1 africa    54
2 americas  35
3 asia     59
4 europe    48
```

```
count(gео, income2017)
```

```
# A tibble: 4 x 2
  income2017      n
  <chr>        <int>
1 high         58
2 low          31
3 lower_mid    52
4 upper_mid    55
```

One annoyance here is that the different categories in `income2017` aren't in a sensible order. This comes up quite often, for example when sorting or plotting categorical data. R's solution is a further type of vector called a *factor* (think a factor of an experimental design). A factor holds categorical data, and has an associated ordered set of *levels*. It is otherwise quite similar to a character vector.

Any sort of vector can be converted to a factor using the `factor` function. This function defaults to placing the levels in alphabetical order, but takes a `levels` argument that can override this.

```
head( factor(geo$income2017, levels=c("low","lower_mid","upper_mid","high")) )
```

```
[1] high      high      lower_mid upper_mid upper_mid high
Levels: low lower_mid upper_mid high
```

We should modify the `income2017` column of the `geo` table in order to use this:

```
geo$income2017 <- factor(geo$income2017, levels=c("low","lower_mid","upper_mid","high"))
```


`count` now produces the desired order of output:

```
count(geo, income2017)
```

```
# A tibble: 4 x 2
  income2017      n
  <fct>        <int>
1 low         31
2 lower_mid   52
3 upper_mid   55
4 high        58
```


When `plot` is given a factor, it shows a bar plot:

```
plot(geo$income2017)
```



When given two factors, it shows a mosaic plot:

```
plot(geo$income2017, factor(geo$oeed))
```



Similarly we can count two categorical columns at once.

```
count(geo, income2017, oeed)
```

```
# A tibble: 6 x 3  
  income2017 oeed      n  
  <fct>      <lgl> <int>  
1 low      FALSE     5  
2 low      TRUE      1  
3 lower_mid FALSE    58  
4 lower_mid TRUE      1  
5 upper_mid FALSE   54  
6 upper_mid TRUE      3  
7 high     FALSE   30  
8 high     TRUE   30
```

1	low	FALSE	31
2	lower_mid	FALSE	52
3	upper_mid	FALSE	53
4	upper_mid	TRUE	2
5	high	FALSE	29
6	high	TRUE	29

2.8 Readability vs tidyness

The counts we obtained counting `income2017` vs `oecd` were properly tidy in the sense of containing a single unit of observation per row. However to view the data, it would be more convenient to have income as columns and OECD membership as rows. We can use the `pivot_wider` function from `tidyr` to achieve this. (This is also sometimes also called a “cast” or a “spread”.)

```
counts <- count(geo, income2017, oecd)
pivot_wider(counts, names_from=income2017, values_from=n)
```

```
# A tibble: 2 x 5
  oecd    low lower_mid upper_mid  high
<lg1> <int>    <int>    <int> <int>
1 FALSE    31      52      53    29
2 TRUE     NA      NA       2    29
```

We could further specify `values_fill=list(n=0)` to fill in the NA values with 0. Or when using `count`, make sure all the relevant columns are factors and specify `.drop=FALSE`.

Tip

Tidying is often the first step when exploring a data-set. The `tidyr`^a package contains a number of useful functions that help tidy (or un-tidy!) data. We’ve just seen `pivot_wider` which spreads two columns into multiple columns. The inverse of `pivot_wider` is `pivot_longer`, which gathers multiple columns into two columns: a column of column names, and a column of values. `pivot_longer` is often the first step when tidying a dataset you have received from the wild. (This is sometimes also called a “melt” or a “gather”.)

Here’s an animation illustrating these functions.^b

^a<http://tidyr.tidyverse.org/>

^b<https://www.garrickadenbuie.com/project/tidyexplain/#pivot-wider-and-longer>

2.9 Sorting

Data frames can be sorted using the `arrange` function in `dplyr`.

```
arrange(g77, lat)
```

```
# A tibble: 196 x 8
  name      region oecd g77    lat long income2017 southern
  <chr>    <chr>   <lgl> <lgl> <dbl> <dbl> <fct>      <lgl>
1 New Zealand asia    TRUE  FALSE -42   174   high      TRUE
2 Argentina  americas FALSE  TRUE  -34   -64   upper_mid TRUE
3 Chile      americas TRUE   TRUE  -33.5 -70.6 high     TRUE
4 Uruguay    americas FALSE  TRUE  -33    -56   high     TRUE
5 Lesotho    africa  FALSE  TRUE  -29.5  28.2 lower_mid TRUE
6 South Africa africa  FALSE  TRUE  -29    24   upper_mid TRUE
7 Swaziland  africa  FALSE  TRUE  -26.5  31.5 lower_mid TRUE
8 Australia  asia    TRUE  FALSE -25   135   high     TRUE
9 Paraguay   americas FALSE  TRUE  -23.3 -58   upper_mid TRUE
10 Botswana  africa  FALSE  TRUE  -22    24   upper_mid TRUE
# i 186 more rows
```

Numeric columns are sorted in numeric order. Character columns will be sorted in alphabetical order. Factor columns are sorted in order of their levels. The `desc` helper function can be used to sort in descending order.

```
arrange(g77, desc(name))
```

```
# A tibble: 196 x 8
  name      region oecd g77    lat long income2017 southern
  <chr>    <chr>   <lgl> <lgl> <dbl> <dbl> <fct>      <lgl>
1 Zimbabwe  africa  FALSE  TRUE  -19   29.8 low        TRUE
2 Zambia    africa  FALSE  TRUE  -14.3 28.5 lower_mid TRUE
3 Yemen     asia    FALSE  TRUE  15.5  47.5 lower_mid FALSE
4 Vietnam   asia    FALSE  TRUE  16.2 108. lower_mid FALSE
5 Venezuela americas FALSE  TRUE   8   -66 upper_mid FALSE
6 Vanuatu    asia    FALSE  TRUE  -16   167 lower_mid TRUE
7 Uzbekistan asia    FALSE  FALSE 41.7  63.8 lower_mid FALSE
8 Uruguay    americas FALSE  TRUE  -33   -56 high       TRUE
9 United States americas TRUE   FALSE 39.8 -98.5 high      FALSE
10 United Kingdom europe  TRUE   FALSE 54.8  -2.70 high      FALSE
# i 186 more rows
```

2.10 Joining data frames

Let's move on to a larger data set. This is from the Gapminder⁴ project and contains information about countries over time.

```
gap <- read_csv("r-intro/gap-minder.csv")
gap
```

```
# A tibble: 4,312 x 5
  name          year population gdp_percap life_exp
  <chr>         <dbl>     <dbl>     <dbl>     <dbl>
1 Afghanistan  1800     3280000     603     28.2
2 Albania      1800     410445      667     35.4
3 Algeria      1800    2503218     715     28.8
4 Andorra      1800       2654    1197      NA
5 Angola       1800    1567028     618     27.0
6 Antigua and Barbuda 1800       37000     757     33.5
7 Argentina    1800     534000    1507     33.2
8 Armenia      1800     413326     514      34
9 Australia    1800     351014     814     34.0
10 Austria     1800    3205587    1847     34.4
# i 4,302 more rows
```

Quiz

What does each row represent in this new data frame?

It would be useful to have general information about countries from `geo` available as columns when we use this data frame. `gap` and `geo` share a column called `name` which can be used to match rows from one to the other.

```
gap_geo <- left_join(gap, geo, by="name")
gap_geo
```

```
# A tibble: 4,312 x 12
  name          year population gdp_percap life_exp region oecd g77    lat    long
  <chr>         <dbl>     <dbl>     <dbl>     <dbl> <chr>  <lgl> <lgl> <dbl> <dbl>
1 Afghani~  1800     3280000     603     28.2 asia   FALSE TRUE   33    66
2 Albania   1800     410445      667     35.4 europe FALSE FALSE  41    20
```

⁴<https://www.gapminder.org>


```

3 Algeria      1800      2503218          715      28.8 africa FALSE TRUE    28    3
4 Andorra      1800         2654          1197      NA  europe FALSE FALSE  42.5  1.52
5 Angola       1800     1567028           618     27.0 africa FALSE TRUE  -12.5  18.5
6 Antigua~     1800      37000           757     33.5 ameri~ FALSE TRUE   17.0 -61.8
7 Argenti~     1800     534000          1507     33.2 ameri~ FALSE TRUE  -34   -64
8 Armenia      1800     413326           514      34  europe FALSE FALSE  40.2  45
9 Austral~     1800     351014           814     34.0 asia   TRUE  FALSE  -25   135
10 Austria     1800     3205587          1847     34.4 europe TRUE   FALSE  47.3  13.3
# i 4,302 more rows
# i 2 more variables: income2017 <fct>, southern <lgl>

```

The output contains all ways of pairing up rows by **name**. In this case each row of **geo** pairs up with multiple rows of **gap**.

Tip

The “left” in “left join” refers to how rows that can’t be paired up are handled. **left_join** keeps all rows from the first data frame but not the second. This is a good default when the intent is to attaching some extra information to a data frame. **inner_join** discards all rows that can’t be paired up. **full_join** keeps all rows from both data frames. Here are some animations illustrating joins.^a

^a<https://www.garrickadenbuie.com/project/tidyexplain/#mutating-joins>

2.11 Further reading

We’ve covered the fundamentals of dplyr and data frames, but there is much more to learn. Notably, we haven’t covered the use of the pipe `|>` to chain dplyr verbs together (and you may also come across an older version of the pipe symbol `%>%`). The “R for Data Science” book⁵ is an excellent source to learn more. The Monash Data Fluency “Programming and Tidy data analysis in R” course⁶ also covers this.

⁵<https://r4ds.hadley.nz/>

⁶<https://monashdatafluency.github.io/r-progtidy/>

3 Plotting with ggplot2

We already saw some of R's built in plotting facilities with the function `plot`. A more recent and much more powerful plotting library is `ggplot2`. `ggplot2` is another mini-language within R, a language for creating plots. It implements ideas from a book called “The Grammar of Graphics”¹. The syntax can be a little strange, but there are plenty of examples in the online documentation².

`ggplot2` is part of the Tidyverse, so loading the `tidyverse` package will load `ggplot2`.

```
library(tidyverse)
```

We continue with the Gapminder dataset, which we loaded with:

```
geo <- read_csv("r-intro/geo.csv")
gap <- read_csv("r-intro/gap-minder.csv")
gap_geo <- left_join(gap, geo, by="name")
```

3.1 Elements of a ggplot

Producing a plot with `ggplot2`, we must give three things:


1. A data frame containing our data.
2. How the columns of the data frame can be translated into positions, colors, sizes, and shapes of graphical elements (“aesthetics”).
3. The actual graphical elements to display (“geometric objects”).

Let's make our first `ggplot`.

```
ggplot(gap_geo, aes(x=year, y=life_exp)) +
  geom_point()
```

¹<https://www.amazon.com/Grammar-Graphics-Statistics-Computing/dp/0387245448>


²<http://ggplot2.tidyverse.org/reference/>



The call to `ggplot` and `aes` sets up the basics of how we are going to represent the various columns of the data frame. `aes` defines the “aesthetics”, which is how columns of the data frame map to graphical attributes such as x and y position, color, size, etc. `aes` is another example of magic “non-standard evaluation”, arguments to `aes` may refer to columns of the data frame directly. We then literally add layers of graphics (“geoms”) to this.

Further aesthetics can be used. Any aesthetic can be either numeric or categorical, an appropriate scale will be used.

```
ggplot(gap_geo, aes(x=year, y=life_exp, color=region, size=population)) +
  geom_point()
```



Challenge: make a ggplot

This R code will get the data from the year 2010:

```
gap2010 <- filter(gap_geo, year == 2010)
```


Create a ggplot of this with:

- `gdp_percap` as x.
- `life_exp` as y.
- `population` as the size.
- `region` as the color.

3.2 Further geoms


To draw lines, we need to use a “group” aesthetic.

```
ggplot(gap_geo, aes(x=year, y=life_exp, group=name, color=region)) +  
  geom_line()
```



A wide variety of geoms are available. Here we show violin plots. Note again the use of the “group” aesthetic, without this ggplot will just show one big violin. `geom_jitter`, `geom_boxplot`, and the `ggbeeswarm`³ package are some further ways to show distributions.

```
ggplot(gap_geo, aes(x=year, y=life_exp, group=year)) +  
  geom_violin()
```




`geom_smooth` can be used to show trends.

³<https://cran.r-project.org/web/packages/ggbeeswarm/index.html>

```
ggplot(gap_geo, aes(x=year, y=life_exp)) +
  geom_point() +
  geom_smooth()
```


`geom_smooth()` using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'



Aesthetics can be specified globally in `ggplot`, or as the first argument to individual geoms. Here, the “group” is applied only to draw the lines, and “color” is used to produce multiple trend lines:

```
ggplot(gap_geo, aes(x=year, y=life_exp)) +
  geom_line(aes(group=name)) +
  geom_smooth(aes(color=oeed))
```

`geom_smooth()` using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'




3.3 Highlighting subsets

Geoms can be added that use a different data frame, using the `data=` argument.

```
gap_australia <- filter(gap_geo, name == "Australia")

ggplot(gap_geo, aes(x=year, y=life_exp, group=name)) +
  geom_line() +
  geom_line(data=gap_australia, color="red", size=2)
```



Notice also that the second `geom_line` has some further arguments controlling its appearance. These are **not** aesthetics, they are not a mapping of data to appearance, but rather a direct specification of the appearance. There isn't an associated scale as when color was an aesthetic.

Discussion: look at the plot

What do you notice about the data before 1870?


A plot can raise questions about how the data was gathered, and informs how you fit models and perform tests on it.

Visualize, visualize, visualize!

3.4 Fine-tuning a plot


Adding `labs` to a ggplot adjusts the labels given to the axes and legends. A plot title can also be specified.

```
ggplot(gap_geo, aes(x=year, y=life_exp)) +  
  geom_point() +  
  labs(x="Year", y="Life expectancy", title="Gapminder")
```



Now the figure has proper labels and titles. We would also like to change the look of it a little, so we apply an alternative theme with `theme_bw()`. There are a variety of themes available in `ggplot2`, and there are packages that define further themes. We would also like the title to be in the center, so we do a further customization with the `theme()` function (for more detail please see the docs `?theme`).


```
ggplot(gap_geo, aes(x=year, y=life_exp)) +  
  geom_point() +  
  labs(x="Year", y="Life expectancy", title="Gapminder") +  
  theme_bw() +  
  theme(plot.title = element_text(hjust = 0.5))
```



Now the figure looks better.

`coord_cartesian` can be used to set the limits of the x and y axes. Suppose we want our y-axis to start at zero.

```
ggplot(gap_geo, aes(x=year, y=life_exp)) +  
  geom_point() +  
  coord_cartesian(ylim=c(0,90))
```



Type `scale_` and press the tab key. You will see functions giving fine-grained controls over various scales (x, y, color, etc). These allow transformations (eg `log10`), and manually specified breaks (labelled values). Very fine grained control is possible over the appearance of ggplots, see the `ggplot2` documentation for details and further examples.

Challenge: refine your ggplot


Continuing with your scatter-plot of the 2010 data, add axis labels to your plot.

Give your x axis a log scale by adding `scale_x_log10()`.

3.5 Faceting

Faceting lets us quickly produce a collection of small plots. The plots all have the same scales and the eye can easily compare them.

```
ggplot(gap_geo, aes(x=year, y=life_exp, group=name)) +
  geom_line() +
  facet_wrap(~ region)
```



Note the use of `~`, which we've not seen before. `~` syntax is used in R to specify dependence on some set of variables, for example when specifying a linear model. Here the information in each plot is dependent on the continent.

Challenge: facet your ggplot

Let's return again to your scatter-plot of the 2010 data.

Adjust your plot to now show data from all years, with each year shown in a separate facet, using `facet_wrap(~ year)`.

Advanced: Highlight Australia in your plot.

3.6 Saving ggplots

The act of plotting a ggplot is actually triggered when it is printed. In an interactive session we are automatically printing each value we calculate, but if you are using it with a programming construct such as a for loop or function you might need to explicitly `print()` the plot.

Ggplots can be saved using `ggsave`.

```
# Plot created but not shown.
p <- ggplot(gap_geo, aes(x=year, y=life_exp)) + geom_point()

# Only when we try to look at the value p is it shown
```

```
p

# Alternatively, we can explicitly print it
print(p)

# To save to a file
ggsave("test.png", p)

# This is an alternative method that works with "base R" plots as well:
png("test.png")
print(p)
dev.off()
```

💡 Tip about sizing

Figures in papers tend to be quite small. This means text must be proportionately larger than we usually show on screen. Dots should also be proportionately larger, and lines proportionately thicker. The way to achieve this using `ggsave` is to specify a small width and height, given in inches. To ensure the output also has good resolution, specify a high dots-per-inch, or use a vector-graphics format such as PDF or SVG.

```
ggsave("test2.png", p, width=3, height=3, dpi=600)
```

4 Summarizing data

Having loaded and thoroughly explored a data set, we are ready to distill it down to concise conclusions. At its simplest, this involves calculating summary statistics like counts, means, and standard deviations. Beyond this is the fitting of models, and hypothesis testing and confidence interval calculation. R has a huge number of packages devoted to these tasks and this is a large part of its appeal, but is beyond the scope of today.

Loading the data as before, if you have not already done so:

```
library(tidyverse)

geo <- read_csv("r-intro/geo.csv")
gap <- read_csv("r-intro/gap-minder.csv")
gap_geo <- left_join(gap, geo, by="name")
```

4.1 Summary functions

R has a variety of functions for summarizing a vector, including: `sum`, `mean`, `min`, `max`, `median`, `sd`.

```
mean( c(1,2,3,4) )
```

```
[1] 2.5
```

We can use these on the Gapminder data.

```
gap2010 <- filter(gap_geo, year == 2010)
sum(gap2010$population)
```

```
[1] 6949495061
```

```
mean(gap2010$life_exp)
```

```
[1] NA
```

4.2 Missing values

Why did `mean` fail? The reason is that `life_exp` contains missing values (NA).

```
gap2010$life_exp
```

```
[1] 56.20 76.31 76.55 82.66 60.08 76.85 75.82 73.34 81.98 80.50 69.13 73.79
[13] 76.03 70.39 76.68 70.43 79.98 71.38 61.82 72.13 71.64 76.75 57.06 74.19
[25] 77.08 73.86 57.89 57.73 66.12 57.25 81.29 72.45 47.48 56.49 79.12 74.59
[37] 76.44 65.93 57.53 60.43 80.40 56.34 76.33 78.39 79.88 77.47 79.49 63.69
[49] 73.04 74.60 76.72 70.52 74.11 60.93 61.66 76.00 61.30 65.28 80.00 81.42
[61] 62.86 65.55 72.82 80.09 62.16 80.41 71.34 71.25 57.99 55.65 65.49 32.11
[73] 71.58 82.61 74.52 82.03 66.20 69.90 74.45 67.24 80.38 81.42 81.69 74.66
[85] 82.85 75.78 68.37 62.76 60.73 70.10 80.13 78.20 68.45 63.80 73.06 79.85
[97] 46.50 60.77 76.10    NA 73.17 81.35 74.01 60.84 53.07 74.46 77.91 59.46
[109] 80.28 63.72 68.23 73.42 75.47 65.38 69.74    NA 66.18 76.36 73.55 54.48
[121] 66.84 58.60    NA 68.26 80.73 80.90 77.36 58.78 60.53 81.04 76.09 65.33
[133]    NA 77.85 58.70 74.07 77.92 69.03 76.30 79.84 79.52 73.66 69.24 64.59
[145]    NA 75.48 71.64 71.46    NA 68.91 75.13 64.01 74.65 73.38 55.05 82.69
[157] 75.52 79.45 61.71 53.13 54.27 81.94 74.42 66.29 70.32 46.98 81.52 82.21
[169] 76.15 79.19 69.61 59.30 76.57 71.10 58.74 69.86 72.56 76.89 78.21 67.94
[181]    NA 56.81 70.41 76.51 80.34 78.74 76.36 68.77 63.02 75.41 72.27 73.07
[193] 67.51 52.02 49.57 58.13
```

R will not ignore these unless we explicitly tell it to with `na.rm=TRUE`.

```
mean(gap2010$life_exp, na.rm=TRUE)
```

```
[1] 70.34005
```

Ideally we should also use `weighted.mean` here, to take population into account.

```
weighted.mean(gap2010$life_exp, gap2010$population, na.rm=TRUE)
```

```
[1] 70.96192
```

NA is a special value. If we try to calculate with NA, the result is NA

```
NA + 1
```

```
[1] NA
```

`is.na` can be used to detect NA values, or `na.omit` can be used to directly remove rows of a data frame containing them.

```
is.na( c(1,2,NA,3) )
```

```
[1] FALSE FALSE  TRUE FALSE
```

```
cleaned <- filter(gap2010, !is.na(life_exp))  
weighted.mean(cleaned$life_exp, cleaned$population)
```

```
[1] 70.96192
```

4.3 Grouped summaries

The `summarize` function in `dplyr` allows summary functions to be applied to data frames.

```
summarize(gap2010, mean_life_exp = weighted.mean(life_exp, population, na.rm=TRUE))
```

```
# A tibble: 1 x 1  
  mean_life_exp  
    <dbl>  
1         71.0
```

So far unremarkable, but `summarize` comes into its own when the `.by` argument is used to group data. (There is also an older style of doing this using a function called `group_by`.)

```
summarize(
  gap_geo,
  mean_life_exp = weighted.mean(life_exp, population, na.rm=TRUE),
  .by = year)
```

```
# A tibble: 22 x 2
  year mean_life_exp
  <dbl>         <dbl>
1  1800          30.9
2  1810          31.1
3  1820          31.2
4  1830          31.4
5  1840          31.4
6  1850          31.6
7  1860          30.3
8  1870          31.5
9  1880          32.0
10 1890          32.5
# i 12 more rows
```

Challenge: summarizing

What is the total population for each year? Plot the result.

Advanced: What is the total GDP for each year? For this you will first need to calculate GDP per capita times the population of each country.

The `.by` argument can be used to group by multiple columns, much like `count`. We need to use `c()` for this, as below. We can use this to see how the rest of the world is catching up to OECD nations in terms of life expectancy.

```
result <- summarize(
  gap_geo,
  mean_life_exp = weighted.mean(life_exp, population, na.rm=TRUE),
  .by = c(year, oecd))
result
```

```
# A tibble: 44 x 3
  year oecd mean_life_exp
  <dbl> <lgl>         <dbl>
1  1800 FALSE          29.9
```




```

2 1800 TRUE      34.7
3 1810 FALSE     29.9
4 1810 TRUE      35.2
5 1820 FALSE     30.0
6 1820 TRUE      35.9
7 1830 FALSE     30.0
8 1830 TRUE      36.2
9 1840 FALSE     30.0
10 1840 TRUE     36.2
# i 34 more rows

```

```
ggplot(result, aes(x=year,y=mean_life_exp,color=oeed)) + geom_line()
```



A similar plot could be produced using `geom_smooth`. Differences here are that we have full control over the summarization process so we were able to use the exact summarization method we want (`weighted.mean` for each year), and we have access to the resulting numeric data as well as the plot. We have reduced a large data set down to a smaller one that distills out one of the stories present in this data. However the earlier visualization and exploration activity using `ggplot2` was essential. It gave us an idea of what sort of variability was present in the data, and any unexpected issues the data might have.

4.4 t-test

We will finish this section by demonstrating a t-test. The main point of this section is to give a flavour of how statistical tests work in R, rather than the details of what a t-test does.

Has life expectancy increased from 2000 to 2010?

```
gap2000 <- filter(gap_geo, year == 2000)
gap2010 <- filter(gap_geo, year == 2010)

t.test(gap2010$life_exp, gap2000$life_exp)
```

Welch Two Sample t-test

```
data: gap2010$life_exp and gap2000$life_exp
t = 3.0341, df = 374.98, p-value = 0.002581
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 1.023455 4.792947
sample estimates:
mean of x mean of y
 70.34005  67.43185
```

Statistical routines often have many ways to tweak the details of their operation. These are specified by further arguments to the function call, to override the default behaviour. By default, `t.test` performs an unpaired t-test, but these are repeated observations of the same countries. We can specify `paired=TRUE` to `t.test` to perform a paired sample t-test and gain some statistical power. Check this by looking at the help page with `?t.test`.

It's important to first check that both data frames are in the same order.

```
all(gap2000$name == gap2010$name)
```

```
[1] TRUE
```


```
t.test(gap2010$life_exp, gap2000$life_exp, paired=TRUE)
```

Paired t-test

```
data: gap2010$life_exp and gap2000$life_exp
t = 13.371, df = 188, p-value < 2.2e-16
alternative hypothesis: true mean difference is not equal to 0
95 percent confidence interval:
 2.479153 3.337249
sample estimates:
mean difference
 2.908201
```

When performing a statistical test, it's good practice to visualize the data to make sure there is nothing funny going on.

```
plot(gap2000$life_exp, gap2010$life_exp)
abline(0,1)
```



This is a visual confirmation of the t-test result. If there were no difference between the years then points would lie approximately evenly above and below the diagonal line, which is clearly not the case. However the outlier may warrant investigation.

5 Thinking in R

The result of a t-test is actually a value we can manipulate further. Two functions help us here. `class` gives the “public face” of a value, and `typeof` gives its underlying type, the way R thinks of it internally. For example numbers are “numeric” and have some representation in computer memory, either “integer” for whole numbers only, or “double” which can hold fractional numbers (stored in memory in a base-2 version of scientific notation).

```
class(42)
```

```
[1] "numeric"
```

```
typeof(42)
```

```
[1] "double"
```

Let’s look at the result of a t-test:

```
result <- t.test(gap2010$life_exp, gap2000$life_exp, paired=TRUE)
class(result)
```

```
[1] "htest"
```

```
typeof(result)
```

```
[1] "list"
```

```
names(result)
```

```
[1] "statistic"  "parameter"  "p.value"    "conf.int"   "estimate"
[6] "null.value" "stderr"     "alternative" "method"     "data.name"
```

```
result$p.value
```

```
[1] 4.301261e-29
```

In R, a t-test is just another function returning just another type of data, so it can also be a building block. The value it returns is a special type of vector called a “list”, but with a public face that presents itself nicely. This is a common pattern in R. Besides printing to the console nicely, this public face may alter the behaviour of generic functions such as `plot` and `summary`.

Similarly a data frame is a list of vectors that is able to present itself nicely.

5.1 Lists

Lists are vectors that can hold anything as elements (even other lists!). It’s possible to create lists with the `list` function. This becomes especially useful once you get into the programming side of R. For example writing your own function that needs to return multiple values, it could do so in the form of a list.

```
mylist <- list(hello=c("Hello","world"), numbers=c(1,2,3,4))  
mylist
```

```
$hello  
[1] "Hello" "world"
```

```
$numbers  
[1] 1 2 3 4
```

```
class(mylist)
```

```
[1] "list"
```

```
typeof(mylist)
```

```
[1] "list"
```

```
names(mylist)
```

```
[1] "hello" "numbers"
```

Accessing lists can be done by name with `$` or by position with `[[]]`.

```
mylist$hello
```

```
[1] "Hello" "world"
```

```
mylist[[2]]
```

```
[1] 1 2 3 4
```

5.2 Other types not covered here

Matrices are another tabular data type. These come up when doing more mathematical tasks in R. They are also commonly used in bioinformatics, for example to represent RNA-Seq count data. A matrix, as compared to a data frame:

- contains only one type of data, usually numeric (rather than different types in different columns).
- commonly has `rownames` as well as `colnames`. (Base R data frames can have `rownames` too, but it is easier to have any unique identifier as a normal column instead.)
- has individual cells as the unit of observation (rather than rows).

Matrices can be created using `as.matrix` from a data frame, `matrix` from a single vector, or using `rbind` or `cbind` with several vectors.

You may also encounter “S4 objects”, especially if you use Bioconductor¹ packages. The syntax for using these is different again, and uses `@` to access elements.

5.3 Programming

Once you have a useful data analysis, you may want to do it again with different data. You may have some task that needs to be done many times over. This is where programming comes in:

- Writing your own functions².

¹<http://bioconductor.org/>

²<http://r4ds.had.co.nz/functions.html>

- For-loops³ to do things multiple times.
- If-statements⁴ to make decisions.

The “R for Data Science” book⁵ is an excellent source to learn more. Monash Data Fluency “Programming and Tidy data analysis in R” course⁶ also covers this.

³<http://r4ds.had.co.nz/iteration.html>

⁴<http://r4ds.had.co.nz/functions.html#conditional-execution>

⁵<http://r4ds.had.co.nz/>

⁶<https://monashdatafluency.github.io/r-progtidy/>

6 Next steps

6.1 Deepen your understanding

Our number one recommendation is to read the book “R for Data Science”¹ by Garrett Golemund and Hadley Wickham.

The R Manuals² are the place to look if you need a precise definition of how R behaves.

6.2 Expand your vocabulary

Have a look at these cheat sheets to see what is possible with R.

- Posit’s collection of cheat sheets³ cover some important newer packages in R.
- An old-school cheat sheet⁴ for dinosaurs and people wishing to go deeper.
- A Bioconductor cheat sheet⁵ for biological data.
- The R Graph Gallery⁶ for visual inspiration.
- The R Graphics Cookbook⁷

6.3 Find packages for specific data types

- CRAN⁸ contains over 20,000 R packages.
- The CRAN task views⁹ provide recommendations for specific topics.
- Bioconductor¹⁰ is another package repository, specifically for working with high-throughput biological data.

¹<https://r4ds.hadley.nz/>

²<https://cran.r-project.org/manuals.html>

³<https://posit.co/resources/cheatsheets/>

⁴<https://cran.r-project.org/doc/contrib/Short-refcard.pdf>

⁵<http://mikelove.github.io/bioc-refcard/>

⁶<https://r-graph-gallery.com/>

⁷<https://r-graphics.org/>

⁸<https://cran.r-project.org/>

⁹<https://cran.r-project.org/web/views/>

¹⁰<https://www.bioconductor.org>

6.4 Some pointers on models and statistics

Statistical tasks such as model fitting, hypothesis testing, confidence interval calculation, and prediction are a large part of R, and one we haven't demonstrated fully today.

For any standard statistical test, there will usually be an R function to perform it. Examples include `t.test` from the previous sections, and also `wilcox.test`, `fisher.test`, `chisq.test`, and `cor.test`. Before applying these functions, you may need to use the methods we've learned today to subset and transform your data, or perform some preliminary summarization such as averaging technical replicates. To make sure there are no problems that might invalidate the results from these tests, always visualize your data. If you are performing many tests, adjust for multiple testing with `p.adjust`.

Going beyond this, linear models and the linear model formula syntax `~` are core to much of what R has to offer statistically. Many statistical techniques take linear models as their starting point, including `limma`¹¹ for differential gene expression, `glm` for logistic regression and generalized linear models, survival analysis with `coxph`, and mixed models to characterize variation within populations. I have developed some workshop material on linear models, available here¹².

- “Statistical Models in S” by J.M. Chambers and T.J. Hastie is the primary reference for this, although there are some small differences between R and its predecessor S.
 - The `emmeans`¹³ package will allow you to sensibly interpret models you obtain. Directly interpreting coefficients in models is sometimes misleading. This package fills in the interpretation step, filling in a missing piece of the original framework.
- “An Introduction to Statistical Learning”¹⁴ by G. James, D. Witten, T. Hastie and R. Tibshirani can be seen as further development of the ideas in “Statistical Models in S”, and is available online. It has more of a machine learning than a statistics flavour to it. (The distinction is fuzzy!)
- “Modern Applied Statistics with S” by W.N. Venables and B.D. Ripley is a well respected reference covering R and S.
- “Linear Models with R” and “Extending the Linear Model with R” by J. Faraway¹⁵ cover linear models, with many practical examples.
- Machine learning is a whole further world of packages...

¹¹<https://bioconductor.org/packages/release/bioc/html/limma.html>

¹²<https://monashdatafluency.github.io/r-linear/>

¹³<https://cran.r-project.org/web/packages/emmeans/index.html>

¹⁴<https://www.statlearning.com/>

¹⁵<https://julianfaraway.github.io/faraway/>

- The Carpentries¹⁶ run workshops on scientific computing and data science topics world-wide. The style of this present workshop is very much based on theirs. Their material is all available on their website.
- Many further resources and tutorials exist online.

¹⁶<https://carpentries.org/>